

# Directory-Based Cache Coherence

# Scalable Shared Memory System

- **Bus-based MP:** good for small multiprocessors, but does **not scale**
  - **Physical constraints**
    - long wires (low clock frequency), arbitration delay
  - **Protocol constraints**
    - Snoopy/broadcast: bandwidth getting **saturated quickly**
  - **Contention everywhere:**
    - bus, snoopers, memory
- **How to scale to larger MP?**

# Scalable Shared Memory System

		<i>Interconnection</i>	
		Bus	Point-to-Point
<i>Protocol</i>	Snoopy	Least scalable	More scalable
	Directory	/	Most scalable

# Ways to Scale Bus-Based MP

- **Approach 1:** replace bus with **point-to-point interconnection**, but keep relying on **snooping/broadcast**
  - e.g., AMD Opteron (mesh), IBM Cell (ring), Intel Larrabee (ring)
  - Still **cannot scale** to hundreds of processors

# Ways to Scale Bus-Based MP

- **Approach 2:** replace broadcast with **directory protocol**
  - e.g. SGI Altix system
  - Leads to **distributed shared memory (DSM)** multiprocessors
  - Scaling to **tens to hundreds** of processors

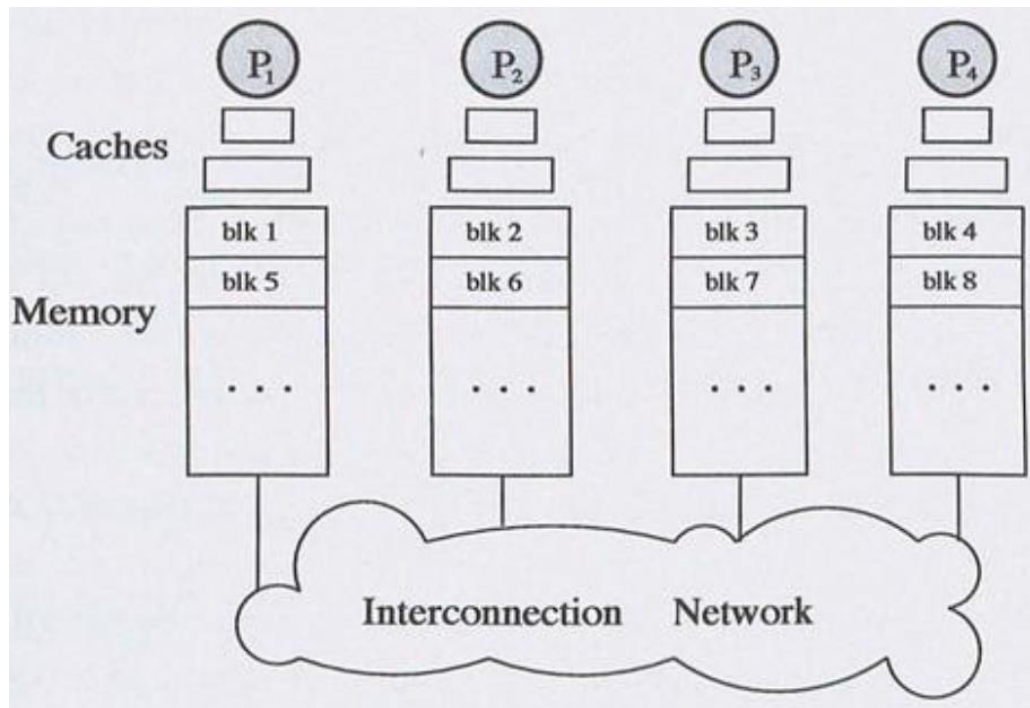
# Directory Protocols

# Questions to be considered

- **Questions:**
  - Which memory a block must be placed in
  - Where the directory of a particular block can be found
  - How the directory information is organized
  - How to handle races in the protocol
  - How to ensure write propagation and write serialization in a DSM
  - .....

# Question 1

- Q1: **which memory to place a particular block?**
- One solution: **block interleaving**

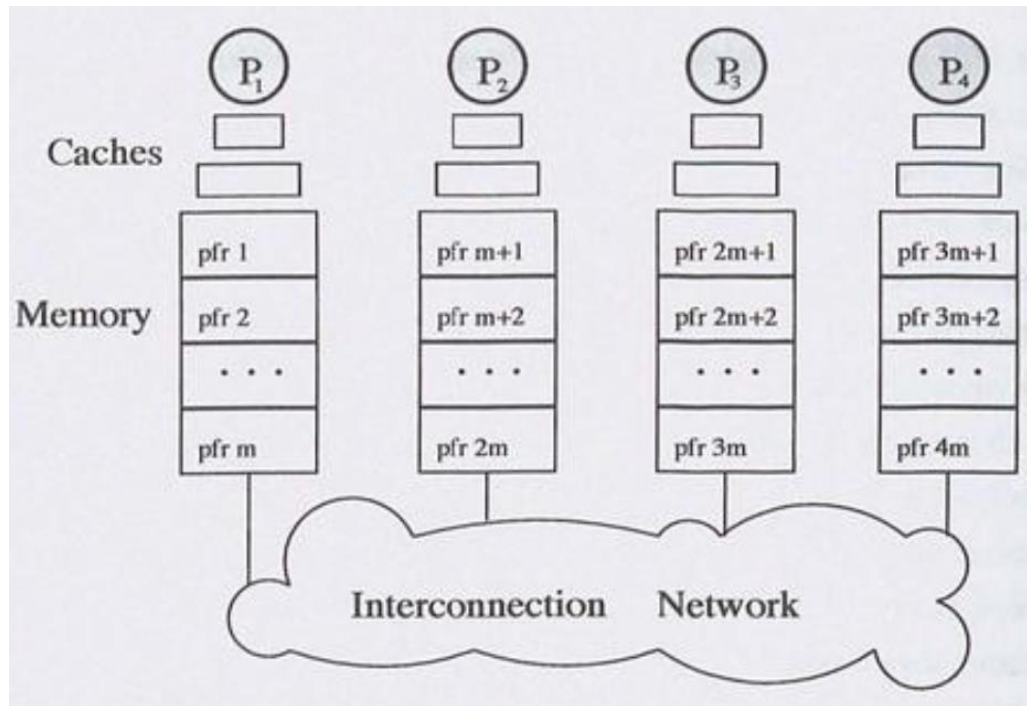


- Works against **spatial locality** of accesses
- In most DSMs, block interleaving is **not used**.



# Question 1

- Another solution: **mapping without interleaving**
  - Consecutive physical page frames (pfr) are mapped to a memory.



- OS must be involved in deciding where to allocate
  - First touch
    - Processor that first accesses it.
  - round robin

# Finding a Block

- **A block can only map to a single memory**
  - This memory is referred to as the “**home**” of the block
- **The home memory keeps track of which caches may have the copy of a block**
  - in a structure called the “**directory**”

# Question 2

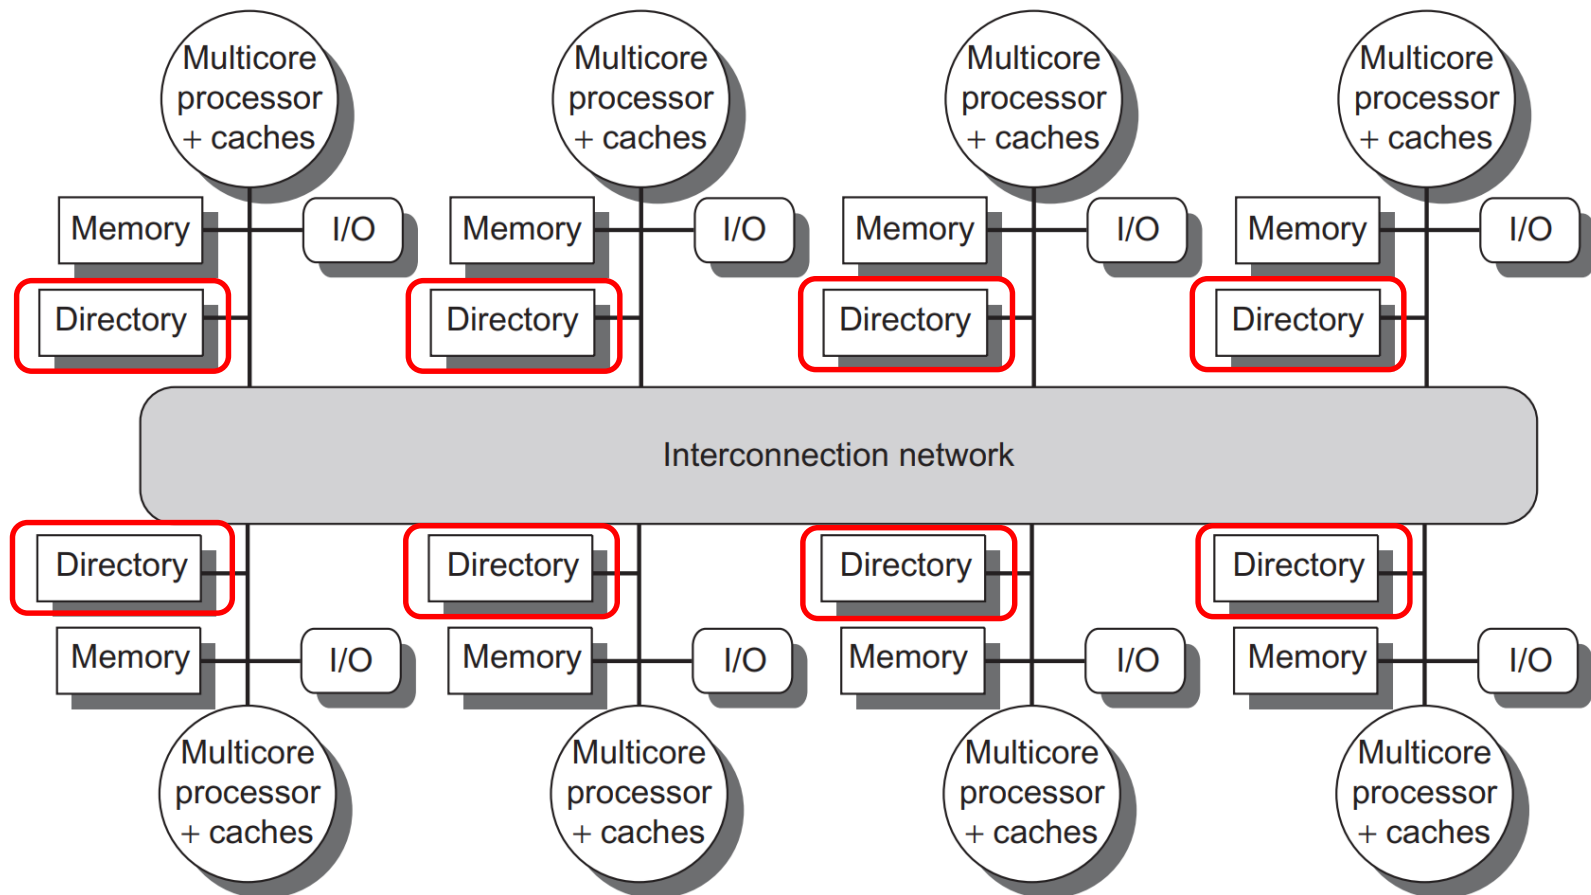
- Q2: **where is directory stored?**
- **Solution:**
  - **Cache-based directory**
    - Kept as a doubly linked list among multiple caches
    - Maintaining the list is *too complicated!*
  - **Memory-based directory**
    - Keep directory in the main memory.
    - **Mainstream** approach

# Centralized Directory


- Use a **centralized directory** to keep track of every block
  - Which caches have each block
  - Dirty status of each block
- ***Not scalable***
  - The directory receives requests from all processors
  - Its central location is far away from most processors

# Distributed Directory

- The directory is scattered at different nodes.
  - Keep directory info of a block in home node




# Basic Directory Handling of Load Request



1. Requestor ( $P_1$ ) sends a Read request to the Home ( $P_2$ ) of the block
2. The home looks up the directory to find out who has the block and in what state. Then it sends an intervention request to the current owner ( $P_3$ )
3. The owner ( $P_3$ ) supplies the block to the requestor ( $P_1$ ) and may update the home as well

# Basic Directory Handling of a Write Req



1. Requestor ( $P_1$ ) sends a Write request to Home ( $P_2$ )
2. Home finds out that  $P_2$  and  $P_3$  are currently sharers, then it sends invalidation messages to them
3. Sharers ( $P_2$  and  $P_3$ ) reply with Invalidation Acknowledgement to  $P_1$

# What Info in the Directory?

- Information that must be kept at the **directory**
  - Which **caches** are currently sharers or owner
  - Which **state** the block is currently cached
- Ideally, if the caches support MESI states, then the **directory keeps MESI states as well**, but:
  - Transition from Exclusive to Modified in a cache is **silent**
  - Hence, the directory **cannot tell** if a block is cached in Exclusive or Modified state



# What Info in the Directory

- Hence, the directory may keep **3 states**
  - **Shared (S)**
    - The block is cached cleanly, possibly by multiple caches
  - **Uncached (U)**
    - The block is uncached or cached in invalid state.
  - **Exclusive or Modified (EM)**
    - The block is either cached in Exclusive or Modified state in **only one cache**.

# Basic DSM Cache Coherence Protocol

# Example

- Cache State: **MESI**
- Directory: **Full-bit Vector**
- 3 States:
  - EM (Exclusive or Modified)
  - S (Shared)
  - U (Uncached)

# Coherence Messages

- Issued by **requesting processor**:
  - **Read**:
    - read request from a processor
  - **ReadX**:
    - **read exclusive (write)** request from a processor that does not have the block
  - **Upgr**:
    - upgrade request (from Shared to Modified), made by a processor that **already has the block**

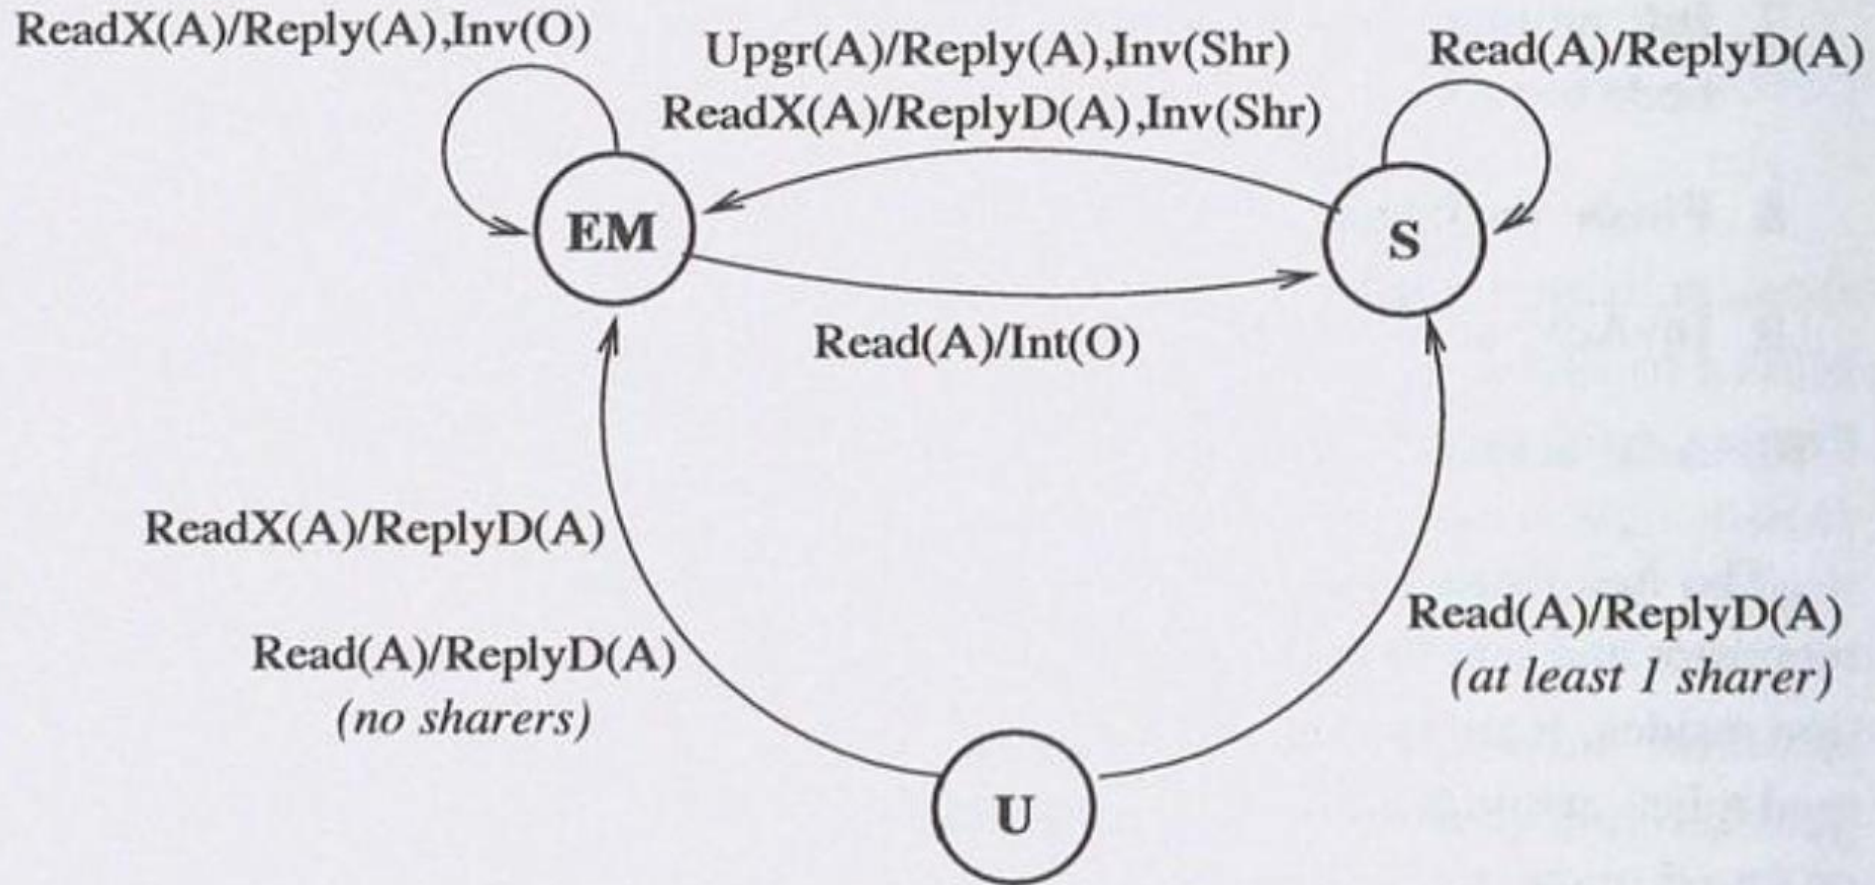
# Coherence Messages

- Issued by **home node**:
  - **ReplyD**:
    - home replies with data to requestor
  - **Reply**:
    - home replies not containing data
  - **Inv (Invalidation)**:
    - Invalidation request sent by home node to sharers
  - **Int (Intervention)**:
    - Intervention request sent by home to owner, asking owner to flush and change to Shared state.

# Coherence Messages

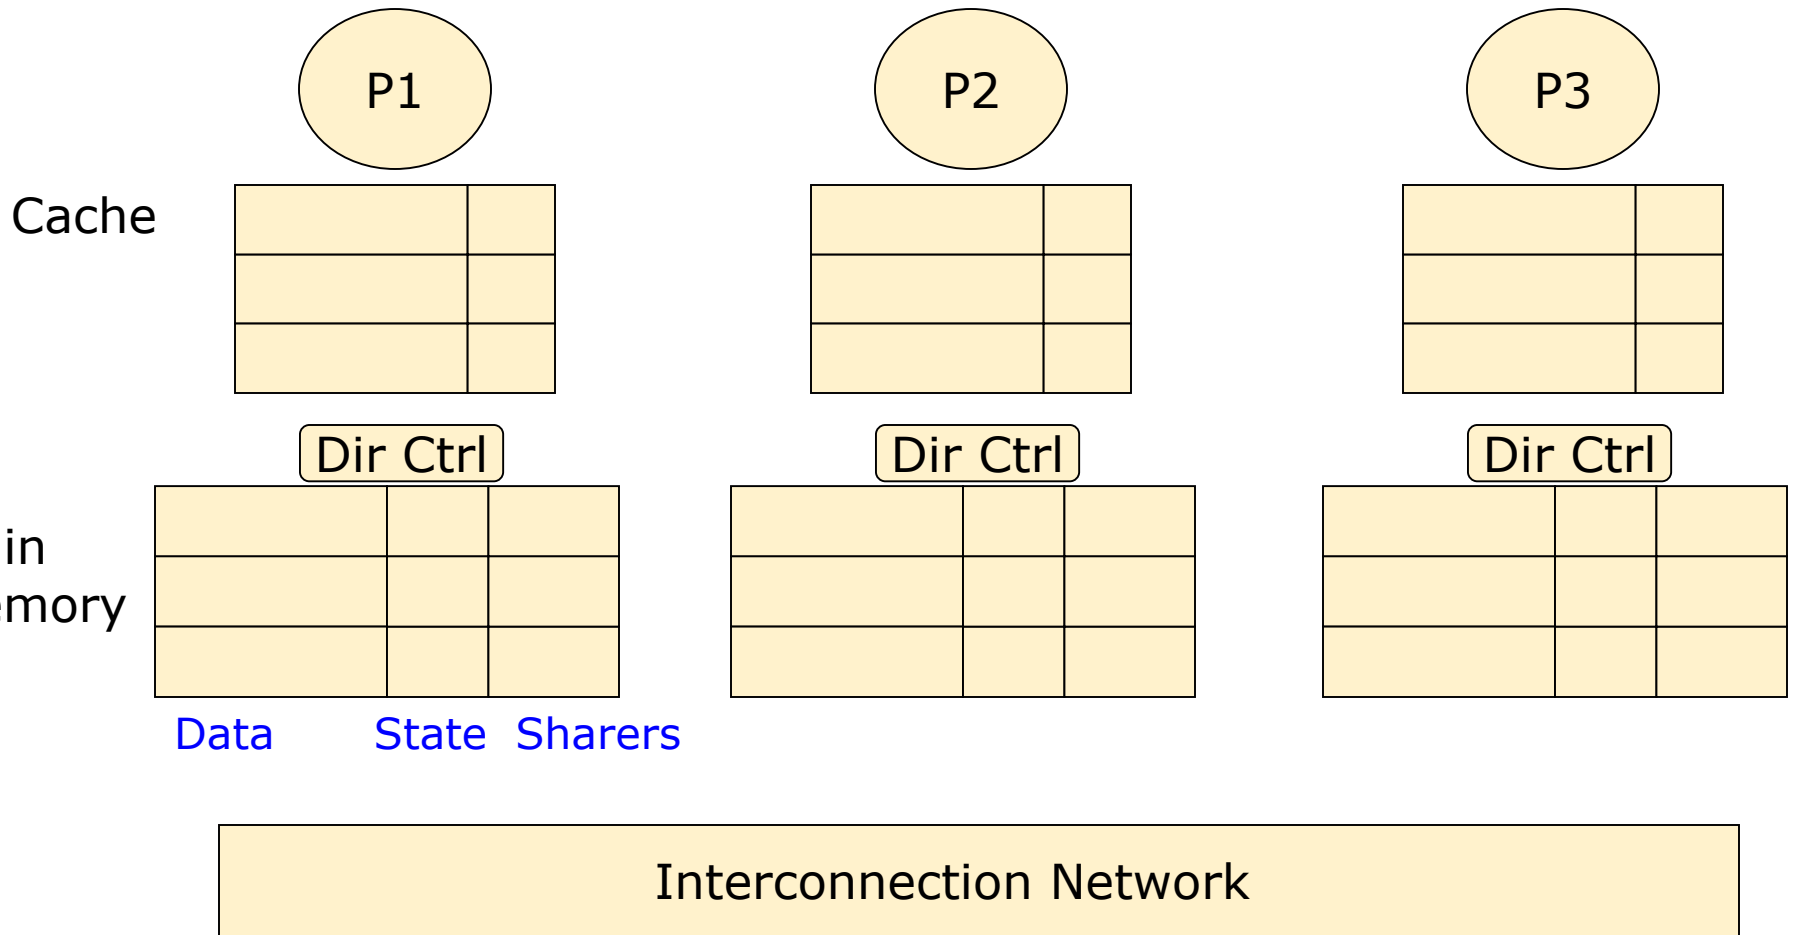
- Issued by **owner of a cache block**:
  - **Flush**:
    - owner flushes data to home + requestor
  - **InvAck**:
    - sharer/owner ack Invalidation msg
  - **Ack**:
    - acknowledgement of receipt of non-invalidation msg

# State Transition at the Directory



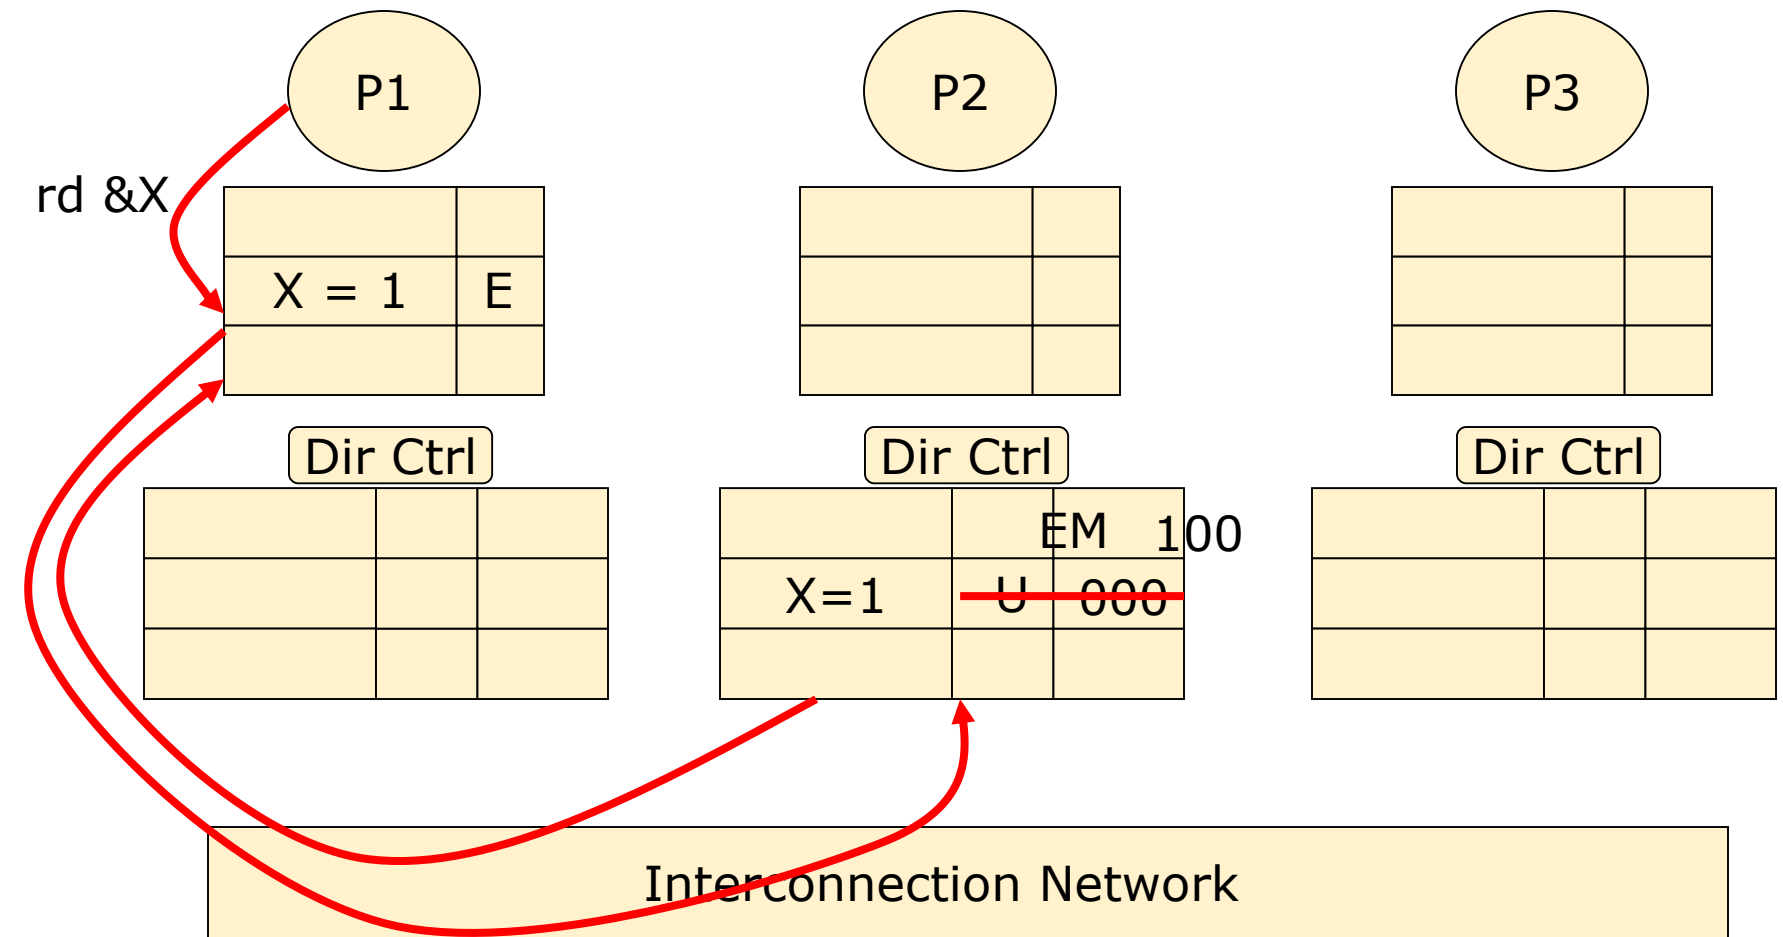
- Legend:
  - <Request-type>(<requestor>) / <Reply-type>(<destination>)
  - O = owner, A = requestor, Shr = sharers

# Full-bit Vector Protocol Visualization

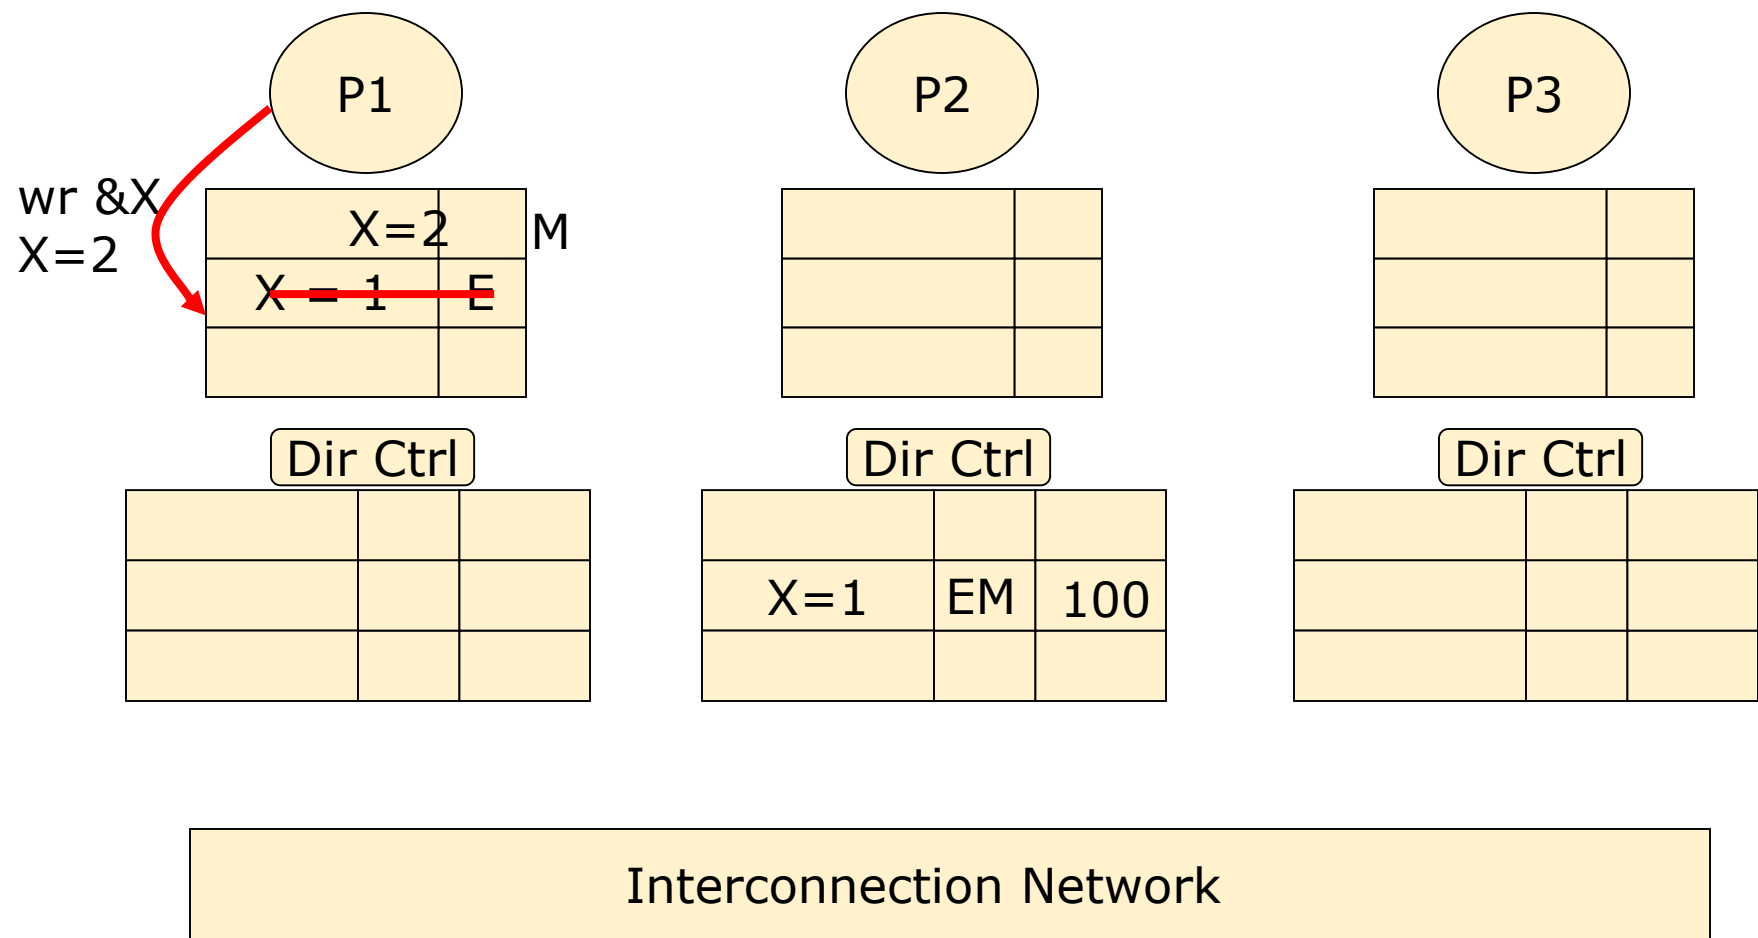




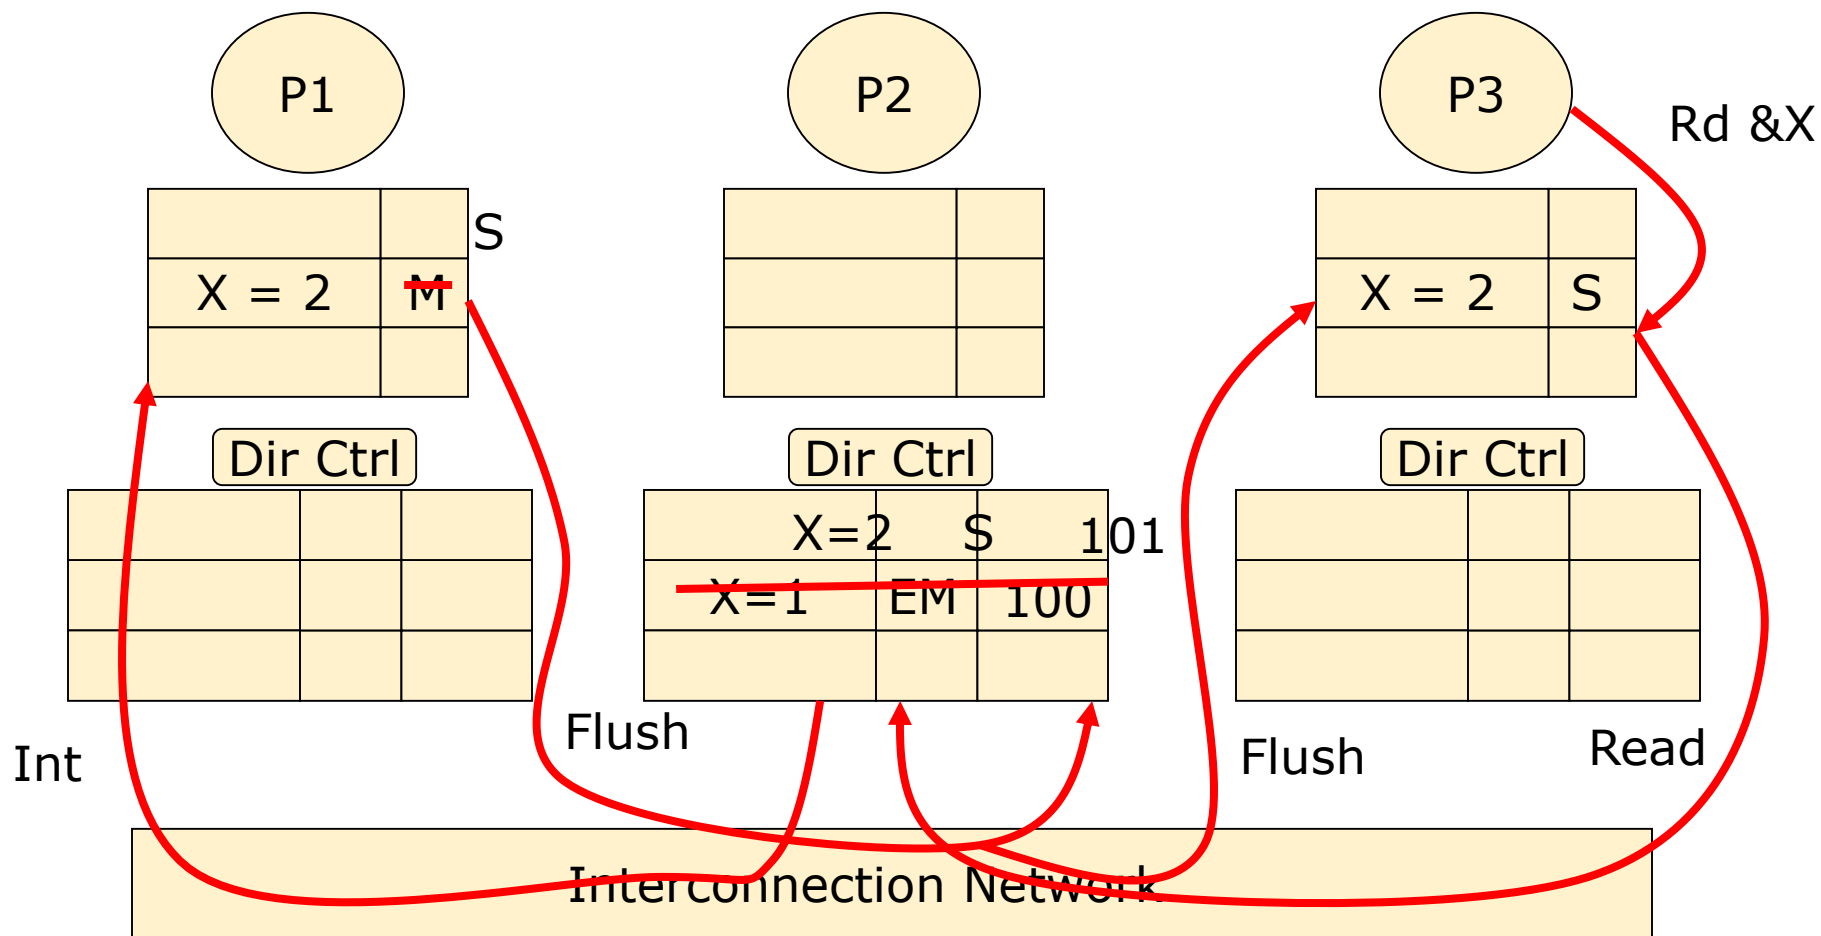
Proc Action	State P1	State P2	State P3	Dir State @Home	Network Msg	Hops
R1	E	-	-	EM, 100	Read (P1-> H), ReplyD (H->P1)	2



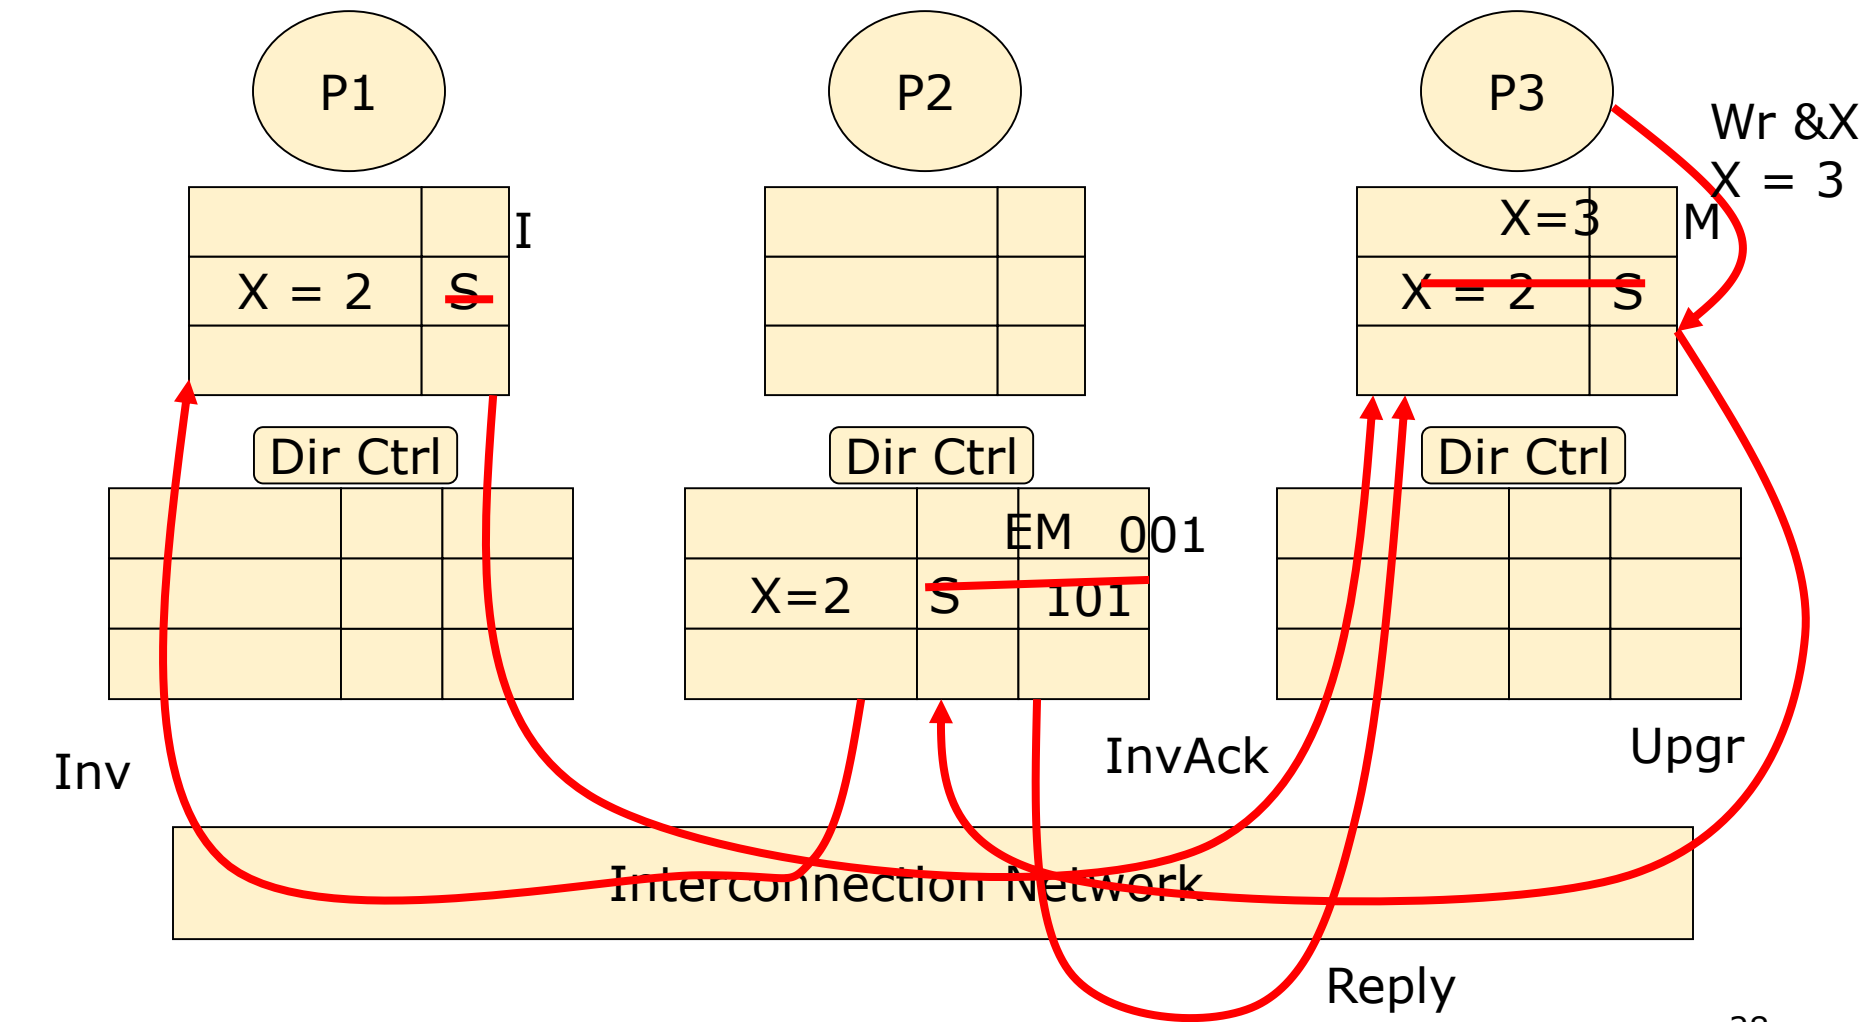
Proc Action	State P1	State P2	State P3	Dir State @Home	Network Msg	Hops
W1	M	-	-	EM, 100	-	0



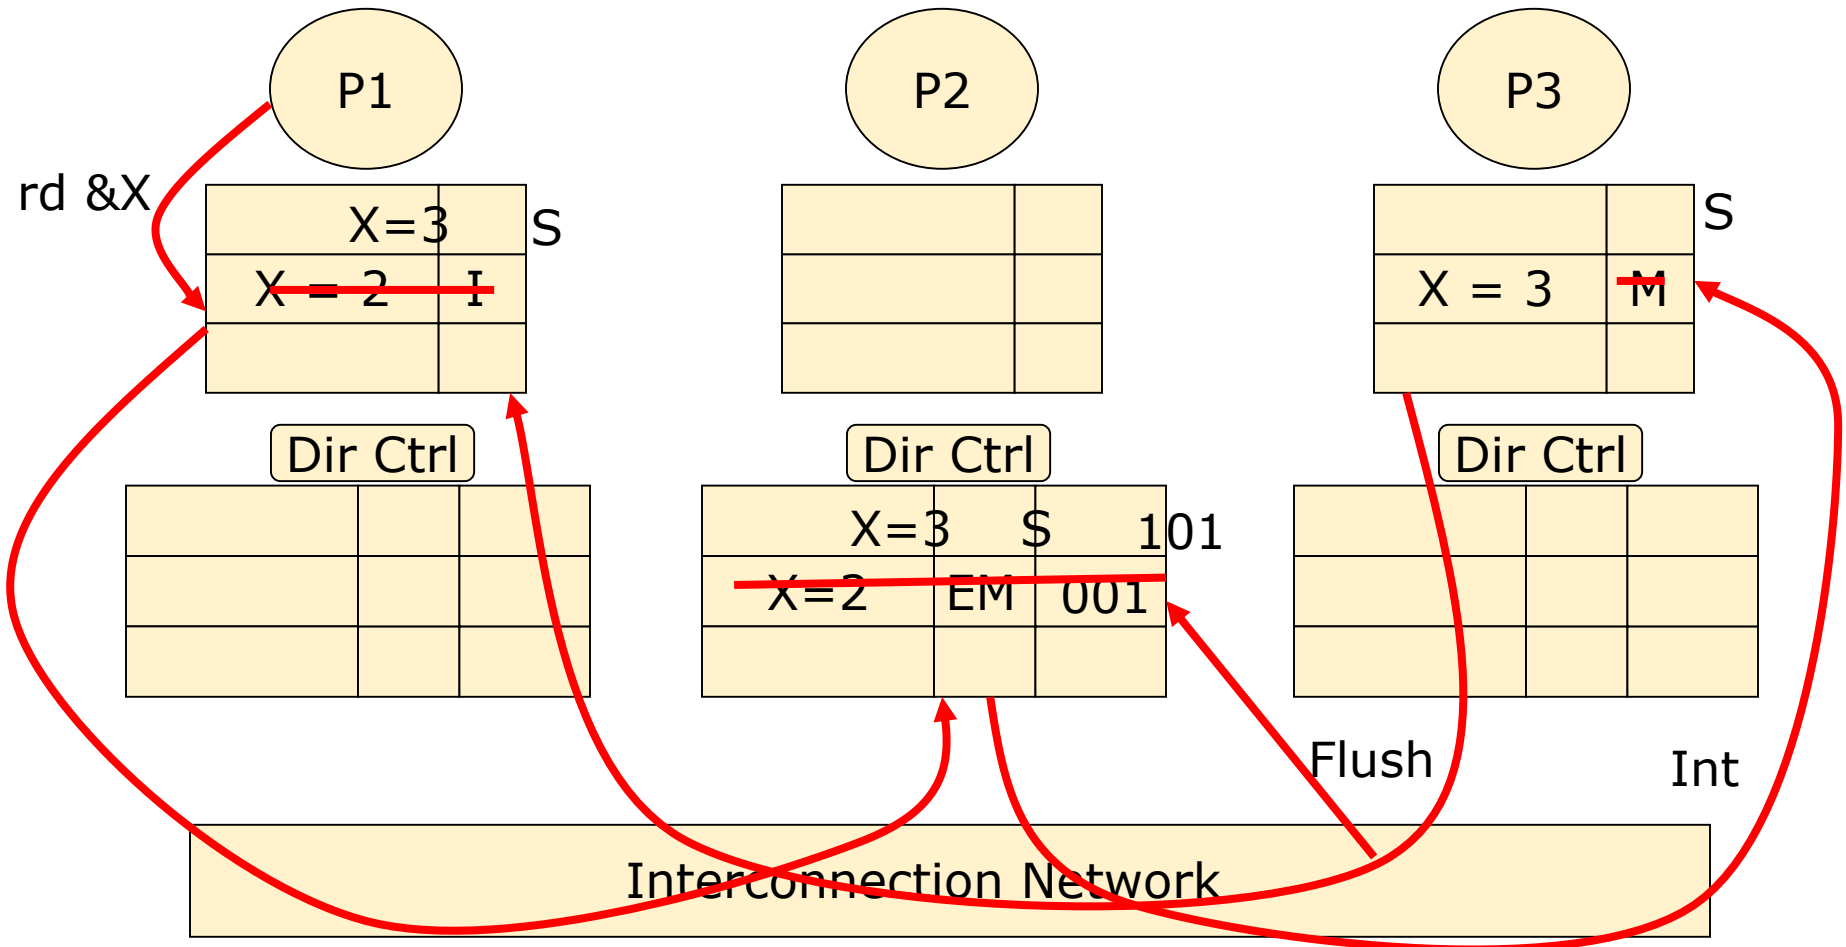
Proc Action	State P1	State P2	State P3	Dir State @Home	Network Msg	Hops
R3	S	-	S	S, 101	Read (P3->H), Int (H->P1), Flush (P1->H, P3)	3



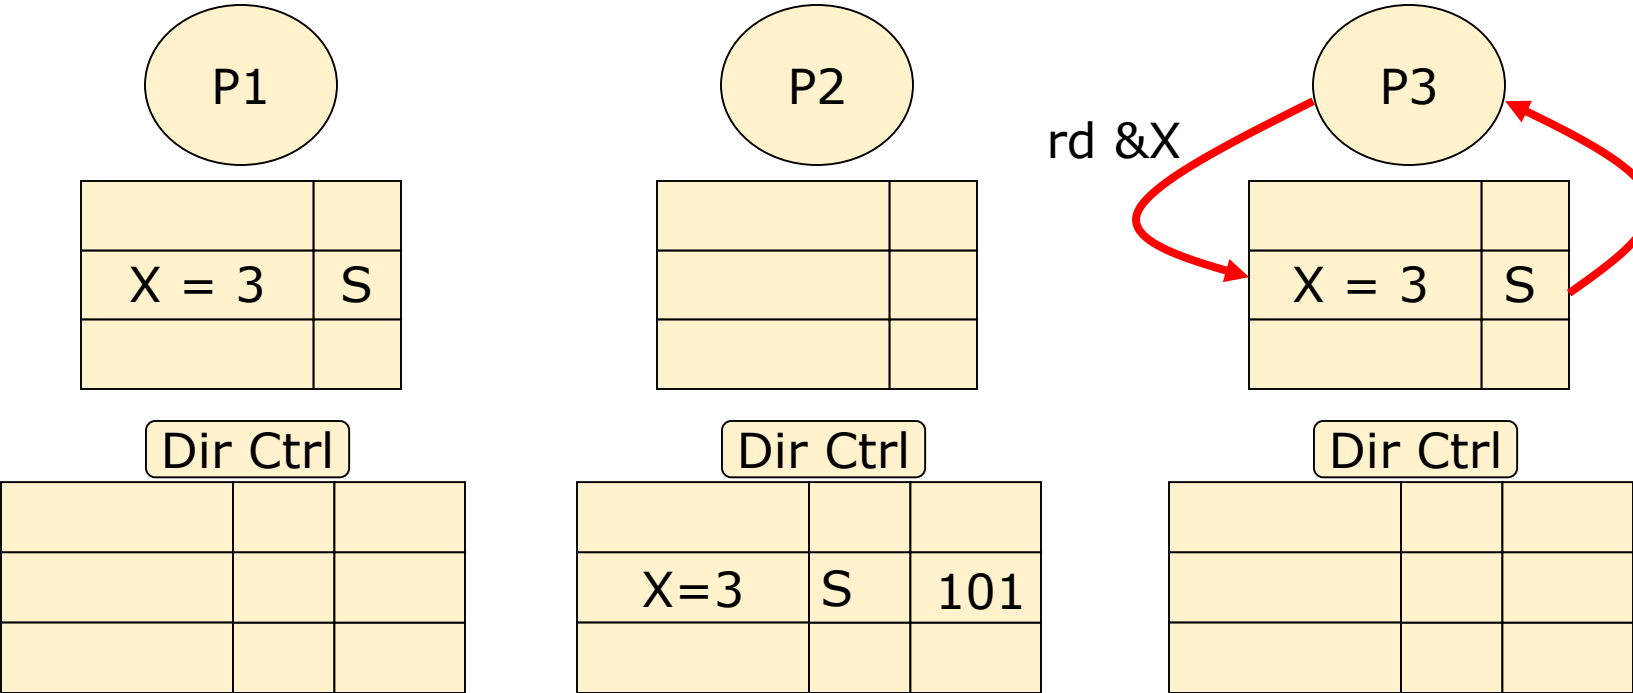
Proc Action	State P1	State P2	State P3	Dir State @Home	Network Msg	Hops
W3	I	-	M	EM, 001	Upgr (P3->H), Reply (H->P3) // Inv (H->P1), InvAck(P1->P3)	3



Proc Action	State P1	State P2	State P3	Dir State @Home	Network Msg	Hops
R1	S	-	S	S, 101	Read (P1->H), Int (H->P3), Flush (P3->H, P1)	3

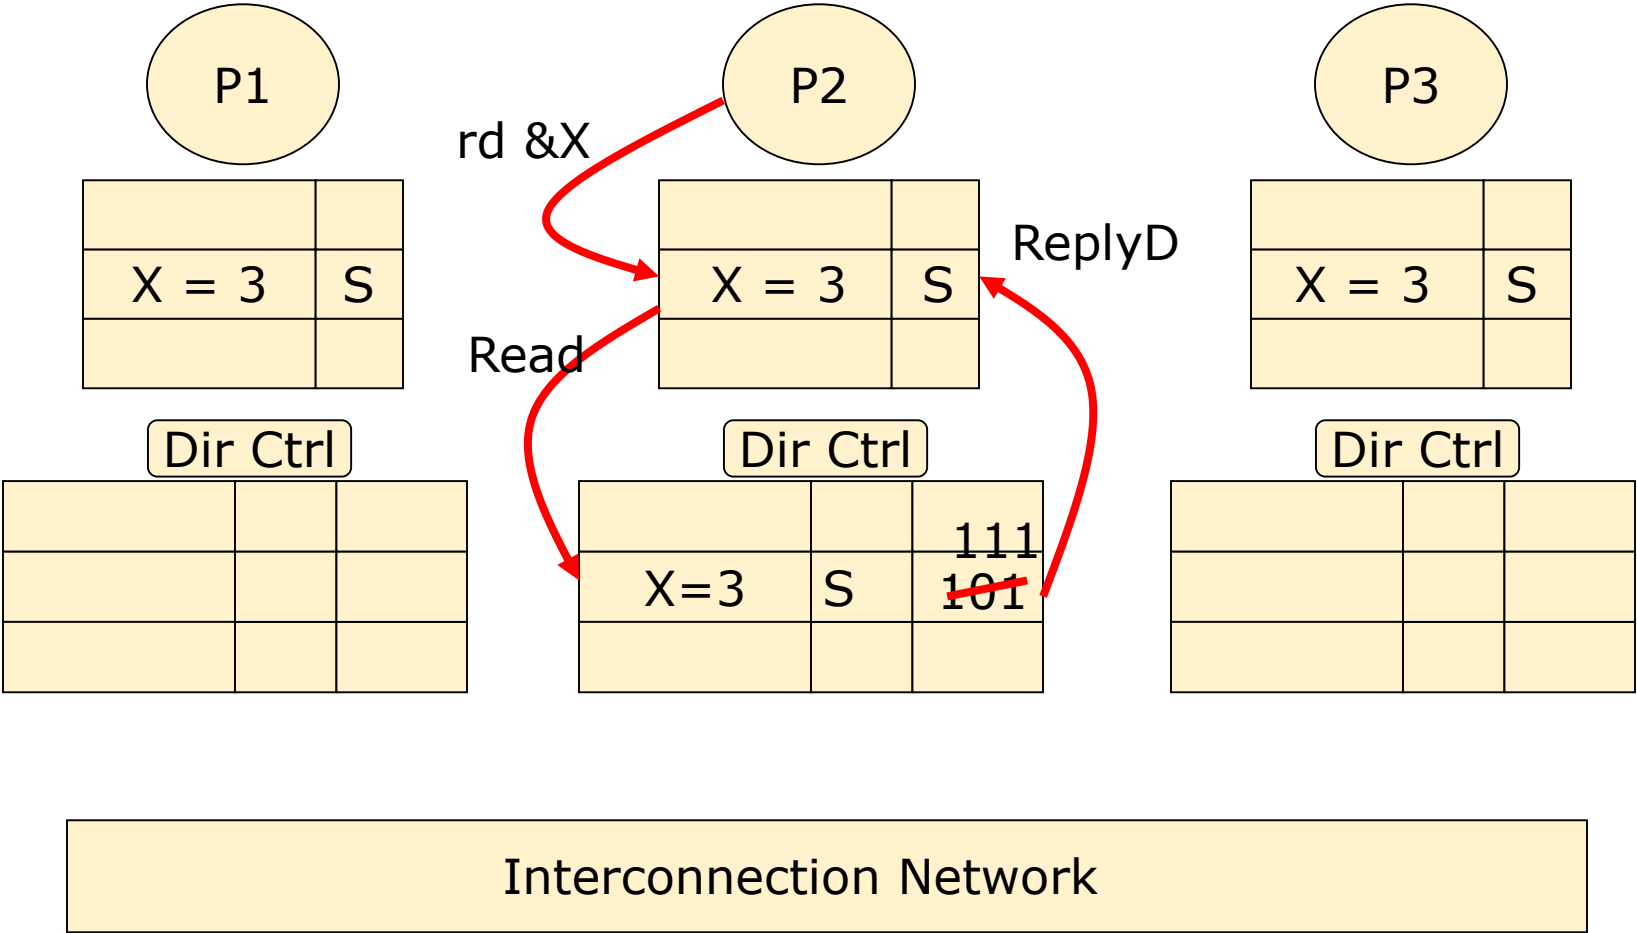


Proc Action	State P1	State P2	State P3	Dir State @Home	Network Msg	Hops
R3	S	-	S	S, 101	-	0



Interconnection Network

Proc Action	State P1	State P2	State P3	Dir State @Home	Network Msg	Hops
R2	S	S	S	S, 111	Read(P2->H), ReplyD(H->P2)	2



# Summary of State Changes

Proc Action	State P1	State P2	State P3	Dir State @Home	Network Msg	Hops
R1	E	-	-	EM, 100	Read (P1-> H), ReplyD (H->P1)	2
W1	M	-	-	EM, 100	-	0
R3	S	-	S	S, 101	Read (P3->H), Int (H->P1), Flush (P1->H, P3)	3
W3	I	-	M	EM, 001	Upgr (P3->H), Reply (H->P3) // Inv (H->P1), InvAck(P1->P3)	3
R1	S	-	S	S, 101	Read (P1->H), Int (H->P3), Flush (P3->H, P1)	3
R3	S	-	S	S, 101	-	0
R2	S	S	S	S, 111	Read(P2->H), ReplyD(H->P2)	2



# Performance Criteria of Directory Schemes

- **Main criteria:**

- Traffic on invalidations
- Latency of invalidations
- Storage overheads

- **Less important criteria:**

- Traffic on interventions
- Latency of interventions

- **Invalidation involves multiple sharers vs. intervention involves a single owner**

**Note:**

Invalidation:

- Any → Invalid

Intervention:

- Exclusive/Modified → Shared

# 课程内容

- 参考材料：
  - Fundamentals of Parallel Computer Architecture.  
Chapter 11 “[Distributed Shared Memory Multiprocessors](#)”.